

# Socket-based Communication

CMPUT-379  
Lab

# Outline

- Inter-process Communication using Sockets
- Unix and I-Net Sockets
- Server side
- Client side
- Server initialization
- Client initialization
- Sending and receiving
- Byte orders
- Closing the connection
- Examples

# IPC using Sockets

- Communication between two processes:
  - Pipes, signals, etc.
  - Sockets
- Sockets provide a communication channel.
- They are essentially file descriptors:
  - We can read/write using regular file commands.
- Processes on the same machine: UNIX Sockets.
- Communication over network: INET Sockets.
- We have a client and a server.

## Socket APIs

- ▶ `socket`: creates a socket of a given domain, type, protocol (buy a phone)
- ▶ `bind`: assigns a name to the socket (get a telephone number)
- ▶ `listen`: specifies the number of pending connections that can be queued for a server socket. (call waiting allowance)
- ▶ `accept`: server accepts a connection request from a client (answer phone)
- ▶ `connect`: client requests a connection request to a server (call)
- ▶ `send`, `sendto`: write to connection (speak)
- ▶ `recv`, `recvfrom`: read from connection (listen)
- ▶ `shutdown`: end the call



# Server side: abstract view

- Server listens for incoming connections.
- Accepts (potentially many) connections.
- What do we need?
  - An "identifier" distinguishing us from other listeners
    - UNIX Communication: a "name".
    - INET Communication: an IP address and port #.
  - A socket on which we accept incoming connections.
  - A way of binding this socket to that "identifier".
  - A way of bringing the server up and listening.
  - A way of accepting new connections.
  - A file descriptor (socket) for the communication.

# Connection-based communication

Server performs the following actions

- ▶ `socket`: create the socket
- ▶ `bind`: give the address of the socket on the server
- ▶ `listen`: specifies the maximum number of connection requests that can be pending for this process
- ▶ `accept`: establish the connection with a specific client
- ▶ `send,recv`: stream-based equivalents of read and write (repeated)
- ▶ `shutdown`: end reading or writing
- ▶ `close`: release kernel data structures

# Client Side: abstract view

- Client initiates the connection to the server.
- What do we need?
  - The server "identifier".
  - A socket to use for the connection.
  - A way of connecting to the server using this socket and that "identifier".

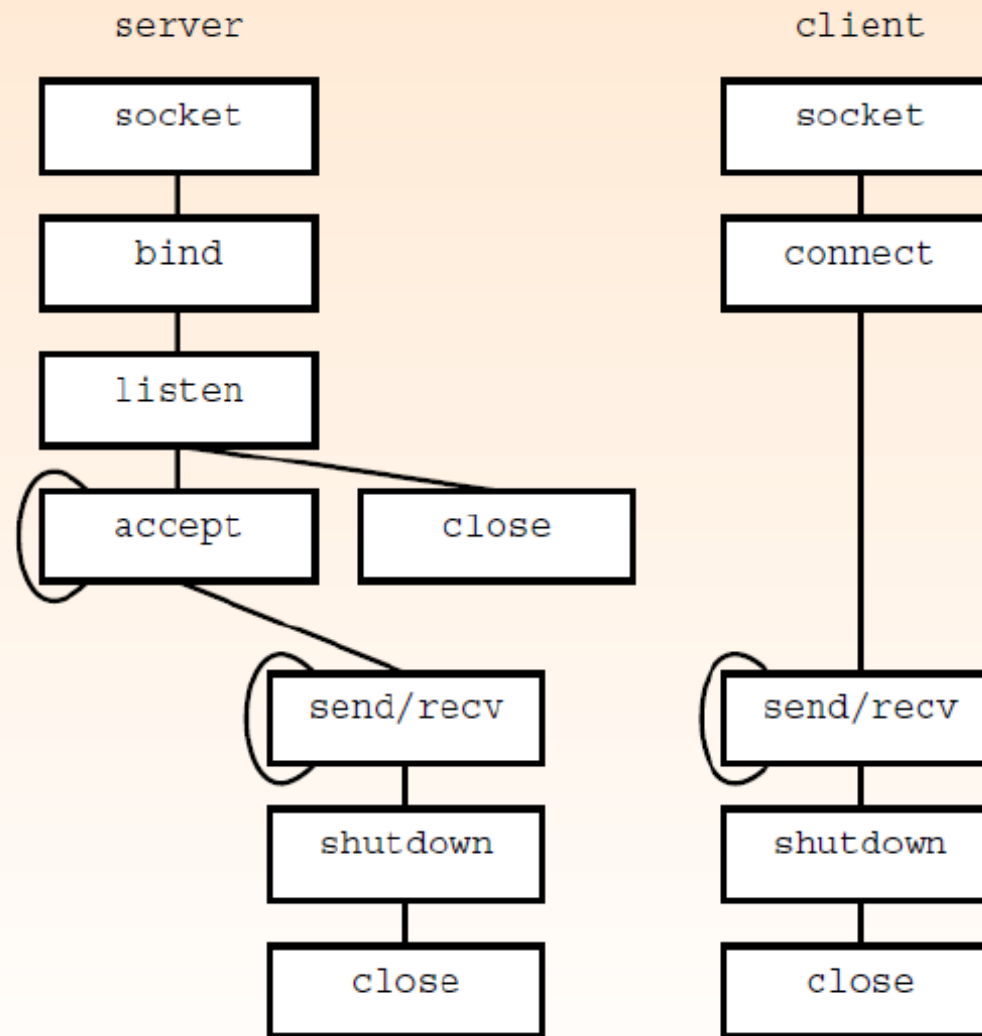
# TCP client

Client performs the following actions

- ▶ `socket`: create the socket
- ▶ `connect`: connect to a server
- ▶ `send,recv`: (repeated)
- ▶ `shutdown`
- ▶ `close`



## TCP-based sockets



# Server Initialization

- Create a socket (click on function names):
  - `s = socket()`
- Set server info (identifier):
  - struct sockaddr\_un
  - struct sockaddr\_in
- Bind the socket to the identifier:
  - bind()
- Start to listen for connections:
  - listen()
- Accept incoming connections:
  - `s = accept()`

# Client Initialization

- Create the socket (click on function names):
  - `s = socket()`
- Set address info (server identifier):
  - `struct sockaddr_un`
  - `struct sockaddr_in`
- Start the connection:
  - `connect()`

# Sending and receiving

- You can use sockets as regular file descriptors.
- To read:
  - read()
- To write
  - write()
- There are also two other functions for sockets:
  - send()
  - recv()
- The difference is that the latter provide socket-specific options and messages.



# Byte orders

- Different machines may have different byte orders
  - Little-Endian vs. Big-Endian.
- Transmitting a number as it is stored in our own machine may lead to incorrect ordering on the other machine. The same for receiving.
- What to do?
  - Have a default "Network Byte Order".
  - Each machine knows his own byte order & network's.
  - Convert all values to and from "Network Byte Order" before sending and after receiving:
    - htons(), ntohs(), htonl(), ntohl()

# Closing the connection

- When you are done, you should close the connection:
  - close()
    - This closes the file descriptor completely.
    - Sending/receiving from the other end generates error.
  - shutdown()
    - This closes parts of a socket functionality.
    - Does not free the file descriptor; you still use `close()`

## Port usage

Note that the initiator of communications needs a fixed port to target communications.

This means that some ports must be reserved for these “well known” ports.

Port usage:

- ▶ 0-1023: These ports can only be binded to by root
- ▶ 1024-5000: well known ports
- ▶ 5001-64K-1: ephemeral ports

# Examples

- Check the four example clients and servers.
- Pay attention to the way the system calls are used.
- Try running them and see the results.
- For in-depth explanation of what you have seen so far, make sure to check the APUE book.



# References

- B. Hall, "Beej's Guide to Network Programming", <http://beej.us/guide/bgnet/>, Accessed: Oct 18, 2011